


# Introduction

- Red-black Tree
  - Self balancing binary search tree with one extra attribute for each node: the colour, which is either red or black.
    - colours - used to ensure that the tree remains balanced during insertions and deletions
  - Although the balance of the tree is not perfect,
    - good enough to reduce the searching time and maintain it around  $O(\log n)$  time
  - Invented in 1972 by Rudolf Bayer

# Introduction

- Red-black Tree Example



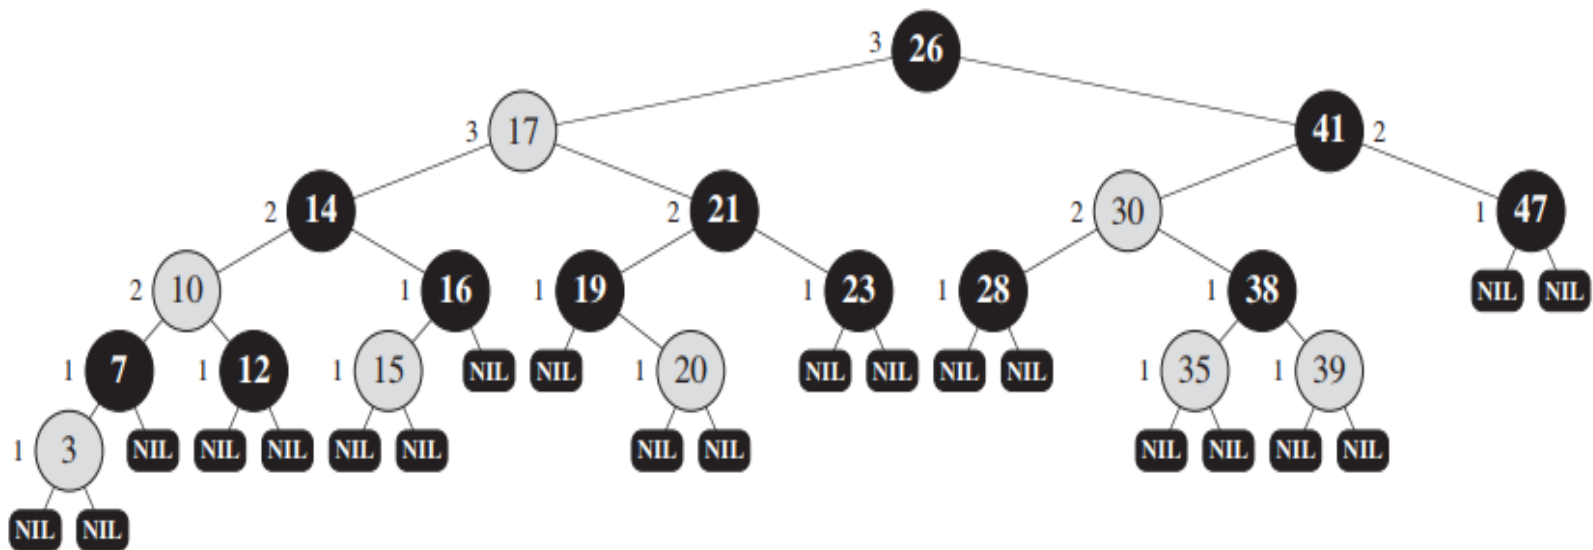
# Red-Black-Trees Properties

- A red-black tree satisfies the following properties:
  1. Every node is either red or black.
  2. The root is black.
  3. Every leaf (NIL) is black.
  4. If a node is red, then both its children are black.
  5. For each node, all simple paths from the node to descendant leaves contain the same number of black nodes.

# Cont...


- Important Note:
  - The longest path from the root node to any leaf node
    - no more than twice as long as the shortest path from the root to any other leaf in that tree
  - Black-height of a node  $x$ :  $bh(x)$ 
    - the number of black nodes (including NIL) on the path from  $x$  to a leaf, not including  $x$

# Cont...




Red-Black Tree

# Cont...



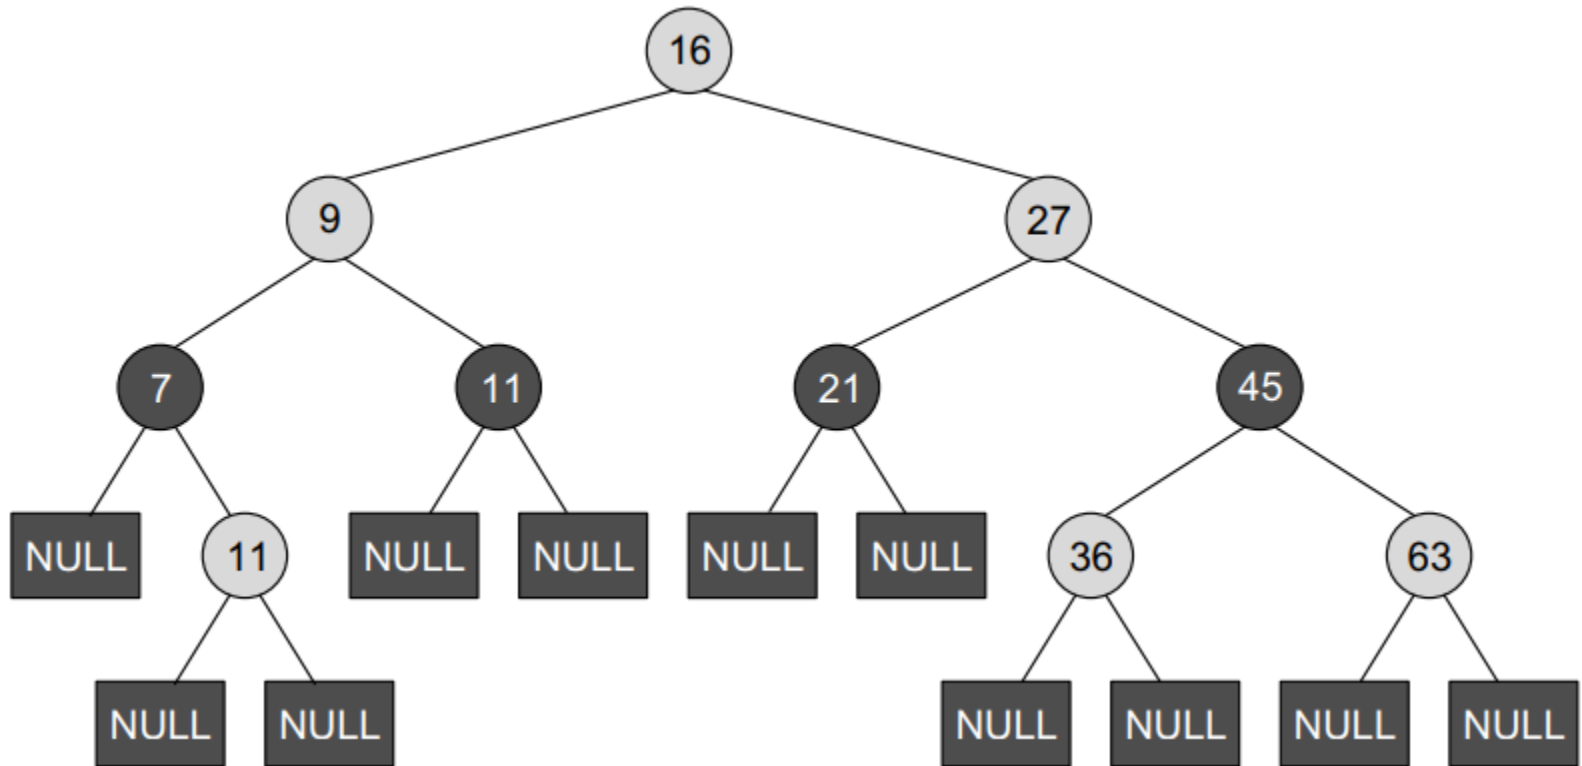
## Red-Black Tree

# Cont...



**Red-Black Tree**

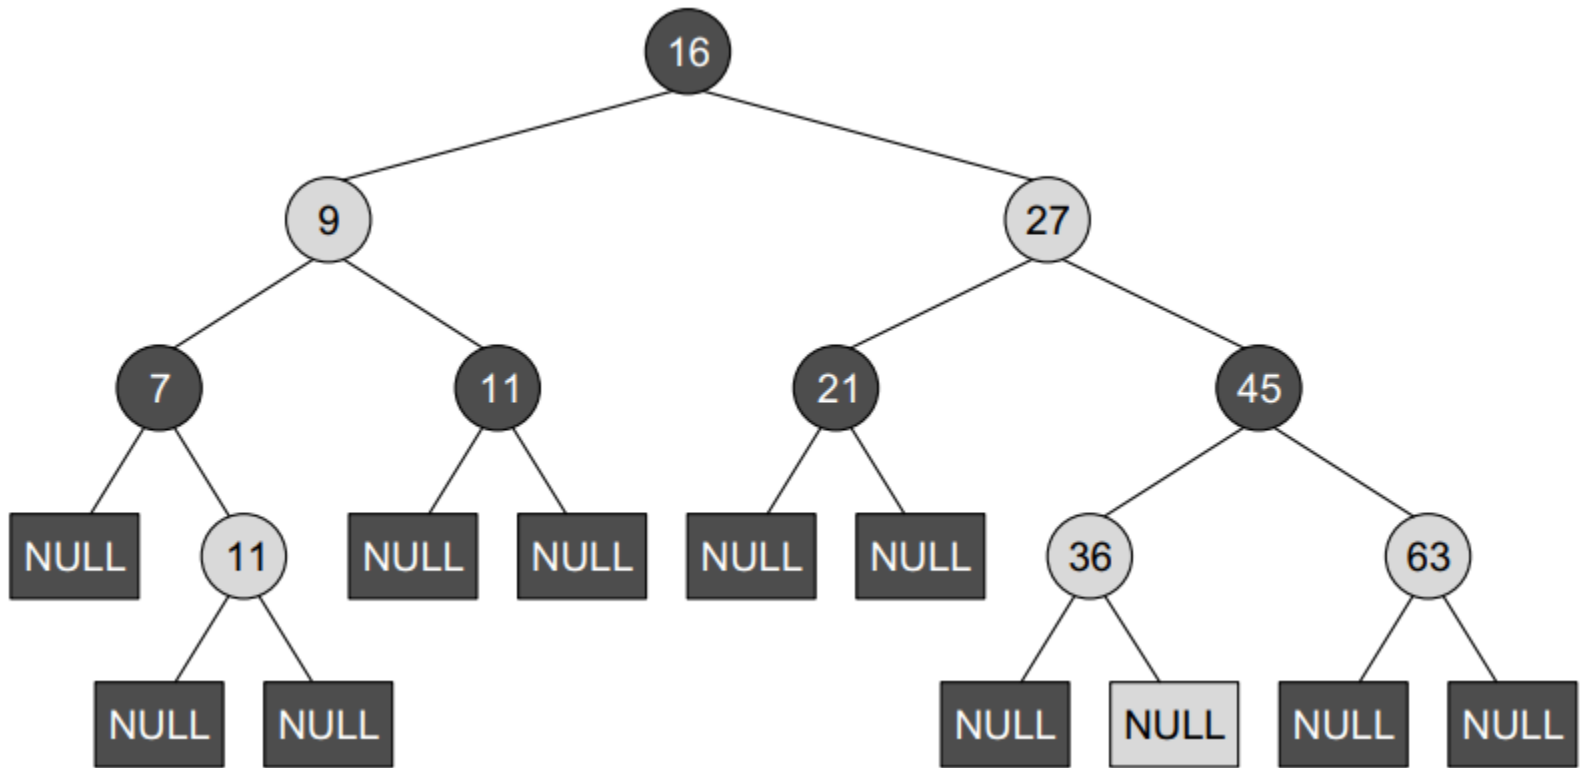
# Cont...



**Not a Red-Black Tree**

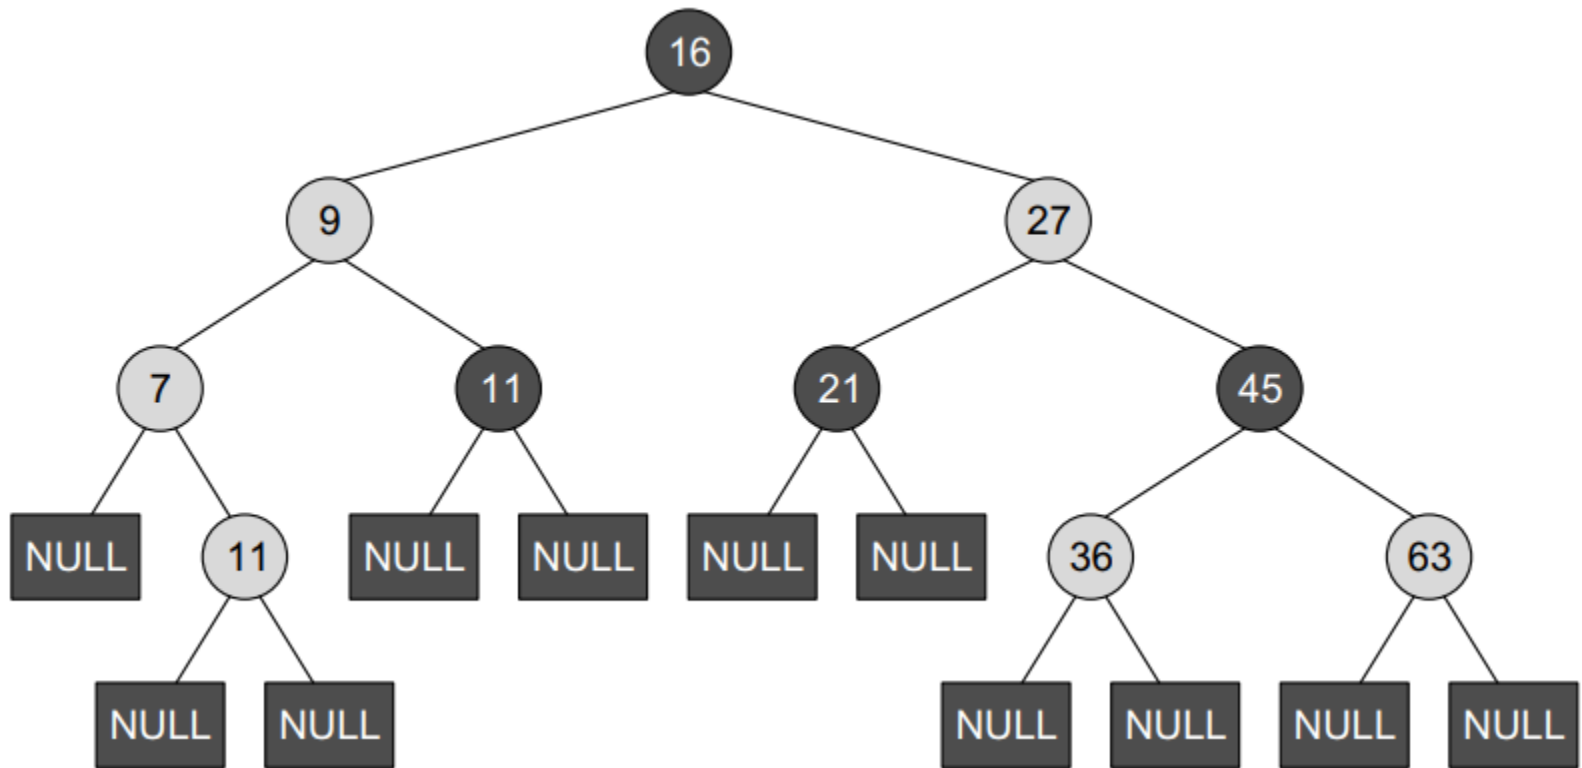


# Cont...



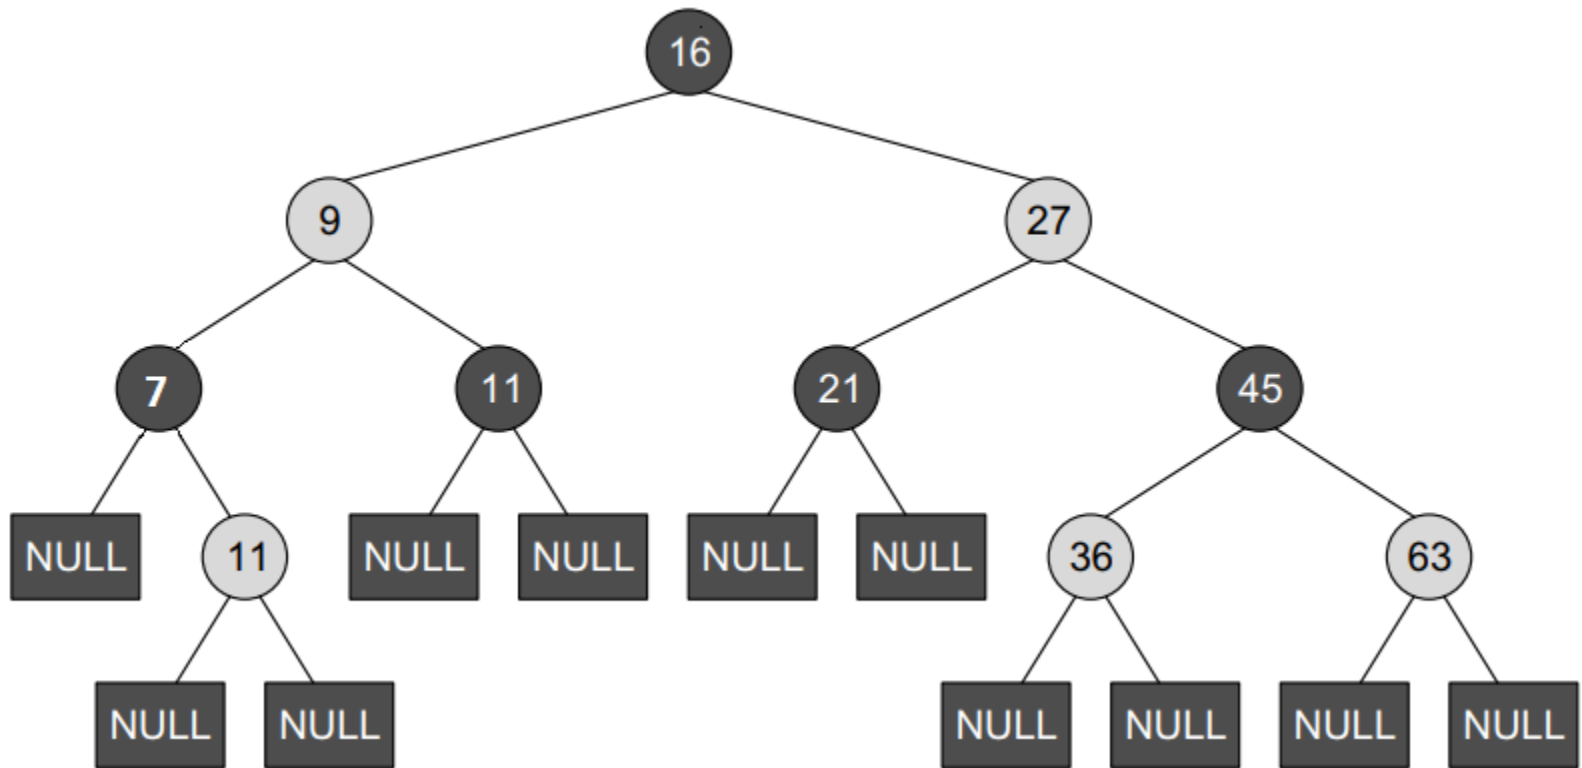
**Not a Red-Black Tree**

# Cont...



**Not a Red-Black Tree**

# Cont...



**Red-Black Tree**

# Cont...

- Why Red-Black Trees?
  - In **BST**, most operations (e.g., search, max, min, insert, delete etc) take  $O(h)$  time, and the cost of these operations may become  $O(n)$  for a skewed Binary tree
  - If the height of the tree remains  $O(\log n)$  after every insertion and deletion, then an upper bound of  $O(\log n)$  for all these operations can be guaranteed
  - In **Red-Black tree**
    - **Height** of a Red-Black tree is always  $O(\log n)$  where  $n$  is the number of nodes in the tree
- How does a Red-Black Tree ensure balance?
  - A chain of 3 nodes is not possible in the Red-Black tree

# Cont...

- Comparison with AVL Tree
  - AVL trees are more balanced compared to Red-Black Trees
  - AVL tree may cause more rotations during insertion and deletion
  - When your application involves frequent insertions and deletions, then Red-Black trees is more preferable

# Cont...


- Interesting points about Red-Black Tree
  - **Black height** of the red-black tree
    - **No. of black nodes** on a path from the **root node** to a **leaf node (excluding the root)**
    - A red-black tree of height  $h$  has **black height  $\geq h/2$**
  - Height of a red-black tree with  $n$  nodes is  **$h \leq 2 \log_2(n+1)$**
  - All leaves (NIL) are black
  - **Black depth** of a node
    - **No. of black nodes** from the **root** to that **node**
      - i.e the number of black ancestors
  - Every red-black tree is a special case of a binary tree

# Cont...

- In the Red-Black tree, the balancing is done with the aid of
  - **Recolouring**
    - the change in colour of the node i.e. if it is red then change it to black and vice versa
  - **Rotation**
    - LL rotation, RR rotation, LR rotation and RL rotation

# Cont...

- Two cases depending upon the colour of the uncle
  - If the uncle is red, we do recolour
  - If the uncle is black, do rotations and/or recolouring






# Insertion in RB Tree

- Insert the node as in a BST and assign a red colour to it
- If the inserted node is a root node
  - change its colour to black
- Else check the colour of the parent node
  - If colour is black then don't change the colour
  - Else colour is red then check colour of the node's uncle
    - If node's uncle has a red colour then change colour of
      - node's parent and uncle to black
      - grandfather to red colour
      - repeat the same process for grandfather


# Cont...

## Uncle is Red

1. Change the colour of X's parent P and uncle U to black.
2. Change the colour of its Grandfather G to red.



Current Tree Structure



Resulting Structure


1. Repeat the all recolouring steps for Grandfather considering it as X.

# Cont...

- If node's **uncle** has a **black** colour then four case are possible
  - Left Left Case (**LL rotation**)
  - Left Right Case (**LR rotation**)
  - Right Right Case (**RR rotation**)
  - Right Left Case (**RL rotation**)

# Cont...


- Left Left Case (RR rotation)



Current Tree Structure

Uncle is Black


1. Right rotation of grandfather G.
2. Then swap the colours of Grandfather G and Parent P.



Resulting Structure


# Cont...

- Right Right Case (LL rotation):



Current Tree Structure

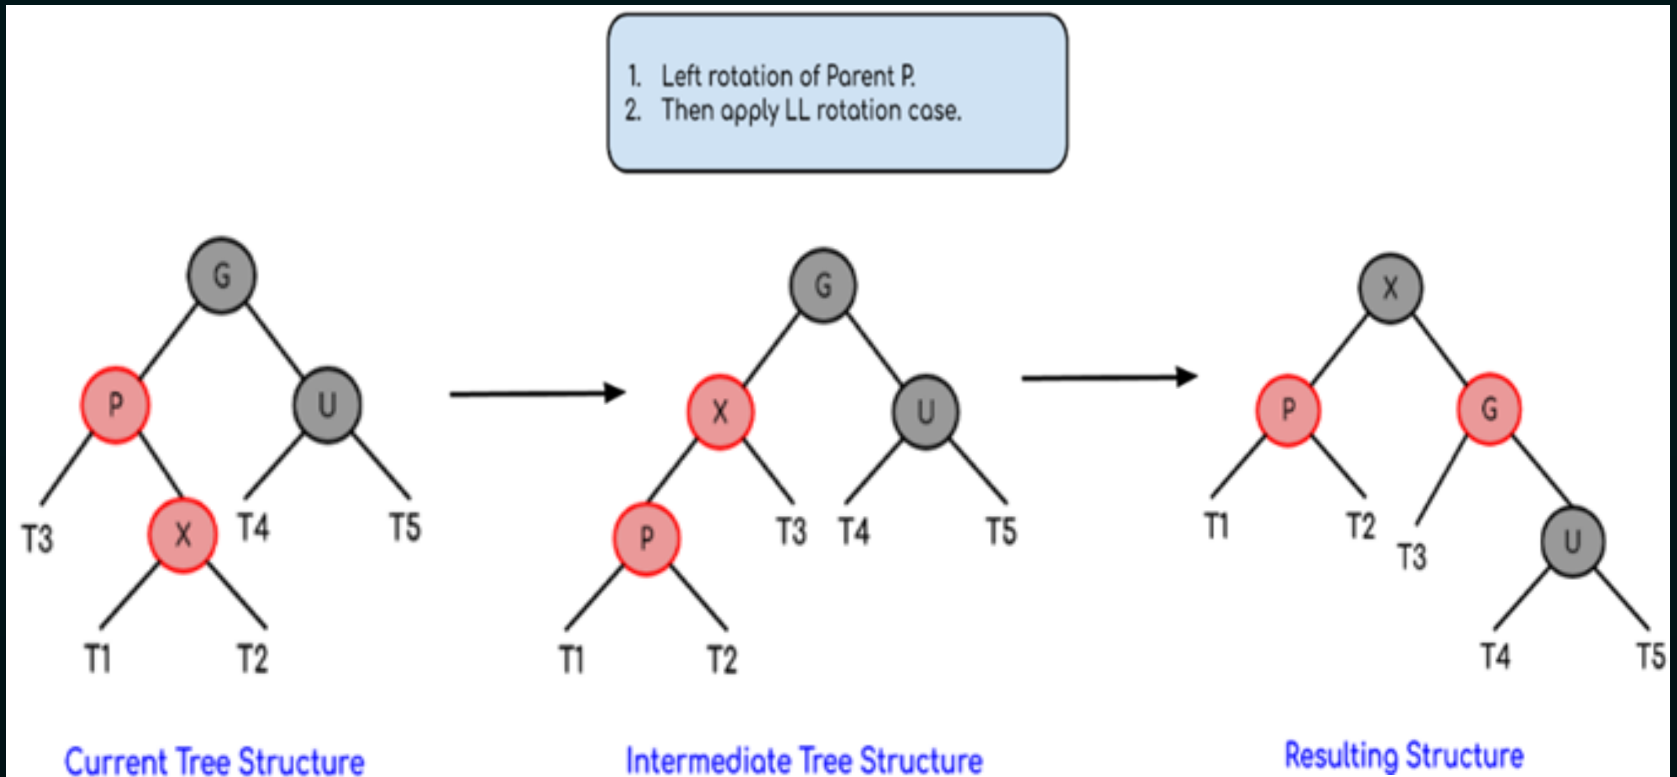
1. Left rotation of grandfather G.
2. Then swap the colours of Grandfather G and Parent P.



Resulting Structure

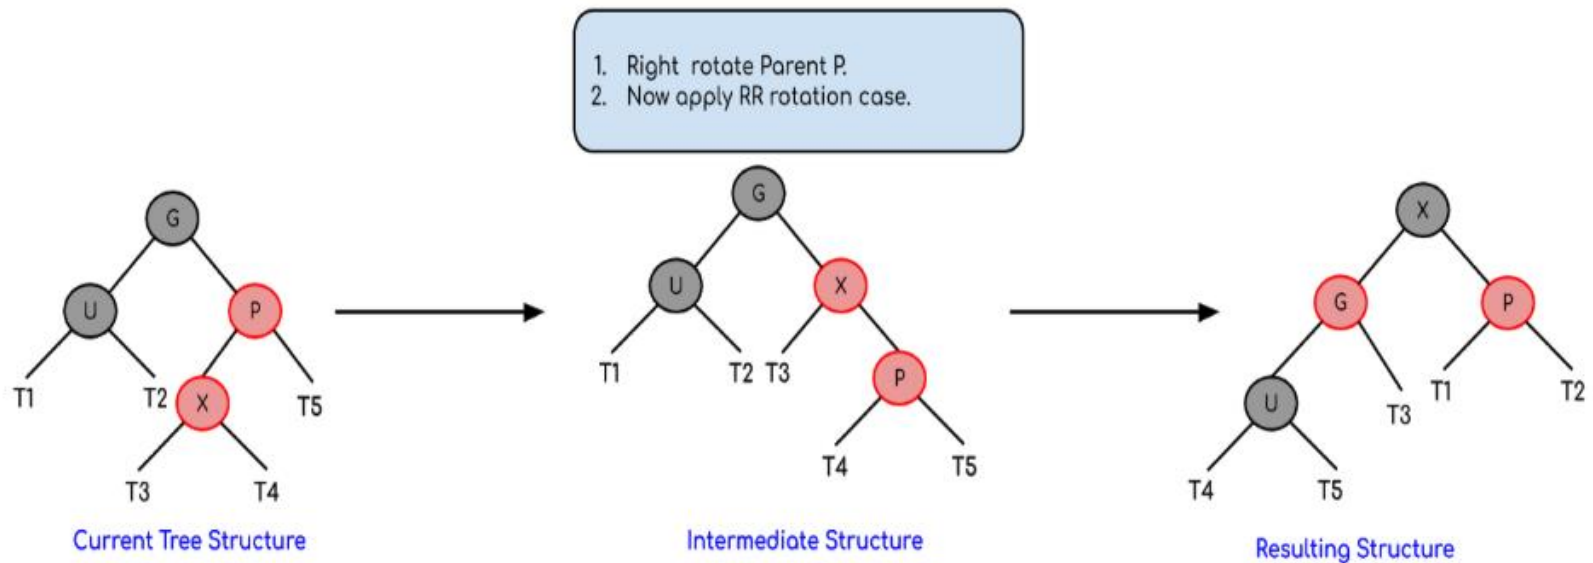
# Cont...

- Left Right Case (LR rotation):



# Cont...

- Right Left Case (RL rotation):



# Cont...

- Insertion Algorithm: x be the newly inserted node
  1. Perform standard BST insertion and make the color of newly inserted nodes as RED
  2. If x is the root, change the color of x as BLACK
  3. x is not the root and the color of x's parent is RED
    - a) If x's uncle is RED (Grandparent must have been black from property 4)
      - i. Change the color of parent and uncle as BLACK.
      - ii. Color of a grandparent as RED.
      - iii. Change x = x's grandparent, repeat steps 2 and 3 for new x.
    - b) If x's uncle is BLACK, then there can be four configurations for x, x's parent (p) and x's grandparent (g)
      - i. Left Left Case (p is left child of g and x is left child of p) –RR rotation
      - ii. Left Right Case (p is left child of g and x is the right child of p) – LR rotation
      - iii. Right Right Case (Mirror of case i) – LL rotation
      - iv. Right Left Case (Mirror of case ii) – RL rotation



# Cont...

- Applications:
  - Most of the self-balancing BST library functions like map and set in C++ (OR TreeSet and TreeMap in Java) use Red-Black Tree.
  - Used to implement CPU Scheduling Linux. Completely Fair Scheduler uses it.
  - Used in the K-mean clustering algorithm for reducing time complexity.
  - MySQL also uses the Red-Black tree for indexes on tables.